

Today: Algorithms for computing max flow

- Recap of the last lecture
- Ford-Fulkerson alg.
- Max Flow Min Cut Theorem
- Maximum bottleneck path & Edmonds-Karp algs.
- Applications of max flow: Bipartite matching problem, (Baseball) Elimination

Last time: Network $G = (V, E, s, t, c)$

(Net) Flow: $f: V \times V \rightarrow \mathbb{R}$ ← defined for u, v ← can be negative

capacities:
 $c: V \times V \rightarrow \mathbb{R}^{\geq 0}$

Convention:
 $c(u, v) = 0 \Rightarrow (u, v) \notin E$

such that:

- $\rightarrow f(u, v) \leq c(u, v) \quad \forall_{u, v}$ (feasibility)
- $\rightarrow \sum_u f(u, v) = 0 \quad \forall_{v \neq s, t}$ (flow consrv.)
- $\rightarrow f(u, v) = -f(v, u) \quad \forall_{u, v}$ (skew symmetry)

Flow value: $|f| = \sum_v f(s, v)$

Maximum Flow problem: Given $G = (V, E, s, t, c)$,

find a flow in G of maximum value

Notation: $F^* = \text{value of max flow (i.e., } |f^*| = F^*)$

max flow
might not
be unique!

s-t cut: A cut $(S, V \setminus S)$ s.t. $s \in S$ & $t \in V \setminus S$

(2)

→ capacity of a cut:

$$c(S) = c(S, V \setminus S) = \sum_{u \in S} \sum_{v \in V \setminus S} c(u, v)$$

(= capacity of edges leaving S)

→ net flow f across the cut:

$$f(S) = f(S, V \setminus S) = \sum_{u \in S} \sum_{v \in V \setminus S} f(u, v)$$

Note:

$$f(S) \leq c(S) \quad \text{for any flow}$$

(*) (by feasibility)

Minimum s-t cut problem: Given $G = (V, E, s, t, c)$,

find an s-t cut of minimum capacity

Claim: For any s-t cut S , and any flow f
(Proved last time)

$$|f| = f(S)$$

⇒ If S^* is a minimum s-t cut and f^* is a max flow,

then

$$F^* = |f^*| = f^*(S^*) \stackrel{(*)}{\leq} c(S^*)$$



Weak duality of flows & s-t cuts
[More in Lecture 9]

(3)

How to find a max flow?Residual network: $G_f = (V, E_f, s, t, c_f)$ of flow f in network G → residual capacities

$$c_f(u, v) = c(u, v) - f(u, v)$$

(= how much extra net $u \rightarrow v$ flow can be directly send)

Note: By feasibility of f , $0 \leq c_f(u, v) \leq c(u, v) + c(v, u)$ → edge $(u, v) \in E_f$ whenever $c_f(u, v) > 0$ (discards saturated edges)

↑
Captures the possibility of pushing some existing $v \rightarrow u$ flow "back".

Augmenting path = directed $s \rightarrow t$ path in G_f → Can be used to push additional flow on such path P up to (residual) bottleneck capacity

$$c_f(P) = \min_{(u, v) \in P} c_f(u, v)$$

(i.e., after the push $|f|$ increases by $c_f(P)$)

Note: $c_f(P) > 0$ since G_f contains only edges with positive c_f

Also: If f is an $s \rightarrow t$ flow in G & f' is an $s \rightarrow t$ flow in G $\Rightarrow f + f'$ is an $s \rightarrow t$ flow in G & $|f + f'| = |f| + |f'|$

Ford-Fulkerson alg. [1956]Idea: Just keep increasing the flow with augmenting pathsAlg.: → $f(u, v) \leftarrow 0 \forall u, v$ (start with zero flow)

This graph changes in each iteration! (as f changes)

Each iteration takes $O(E)$ time (via DFS) → While an augmenting path P exists in G_f : Augment f along P (increasing the value by $c_f(P)$) → Output f



Example execution of F-F alg.:

(4)

Flow:



Residual network:

1)





$$P = (s, v, t) \quad c_p(p) = 2$$

2)





$$P = (s, u, t) \quad c_p(p) = 1$$

3)



4)

$$P = (s, u, v, t) \quad c_p(p) = 1$$



No more augmenting paths!

→ Is this flow a maximum flow?

Yes, as the s-t cut $\{s\}$ has capacity 4

⇒ (by weak duality) $4 \geq F^* = |P| = 4$

F-F alg. works correctly here, but is it correct always?

Max Flow Min Cut Theorem

5

The following statements are equivalent:

- ① $|f| = c(S)$ for some s - t cut S
- ② f is a max flow (i.e., $F^* = |f|$)
- ③ f admits no augmenting path (i.e., there is no s - t path in G_f)

Note: \rightarrow ③ \Rightarrow ② implies that F-F alg. is always correct

\rightarrow ① \Leftrightarrow ② implies that if S^* is the min s - t cut then

$$F^* = c(S^*)$$

Strong duality
of flows & s - t cuts \nearrow
(implies weak duality)

\Rightarrow s - t cuts are the only bottleneck for pushing flows in networks

Proof of the MFMC Thm:

Will show ① \Rightarrow ② \Rightarrow ③ \Rightarrow ①

① \Rightarrow ②: By weak duality, we know that

$$|f| \leq F^* \leq c(S^*) \leq c(S) = |f|$$

$\Rightarrow |f| = F^*$
 \uparrow def. of F^* \nwarrow weak duality \nwarrow def. of S^* \uparrow ①

② \Rightarrow ③: Will show Not ③ \Rightarrow Not ②

But: Not ③ = \exists augmenting path P in G_f

\Rightarrow Augmenting f with P increases $|f|$

$\Rightarrow f$ could not be max. \Rightarrow Not ②

③ \Rightarrow ①: No augmenting path

⑥

\Rightarrow t not reachable from s in G_f

\Rightarrow If $\hat{S} = \{v \mid \exists s \leadsto v \text{ path in } G_f\}$ then
 $\rightarrow s \in \hat{S}$ & $t \notin \hat{S}$

$\Rightarrow \hat{S}$ is an s - t cut & the residual capacity

$$c_f(\hat{S}) = \sum_{u \in \hat{S}} \sum_{v \in V \setminus \hat{S}} c_f(u, v) = 0$$

$$\Rightarrow c_f(\hat{S}) = 0$$

(otherwise $V \setminus \hat{S}$ reachable from \hat{S} (and thus s) in G_f)

$$\Rightarrow \forall_{u \in \hat{S}} \forall_{v \in V \setminus \hat{S}} 0 = c_f(u, v) = c(u, v) - f(u, v)$$

$$\Rightarrow \forall_{u \in \hat{S}} \forall_{v \in V \setminus \hat{S}} c(u, v) = f(u, v)$$

$$\Rightarrow c(\hat{S}) = f(\hat{S}) \Rightarrow \text{①}$$

④

Running time of Ford-Fulkerson alg.?

Recall: Each iteration/augmentation takes $O(E)$ time

How many iterations/augmentations is there?
(in the worst case)

① if capacities are integers $\in [0, C]$ then


$$\# \text{ of augment.} \leq F^* \leq c(\{s\}) \leq |V| \cdot C$$

(since $c_f(p) \geq 1$ always)

see (+) for the reason why

($\leq |V|$ edges leaving s)

\rightarrow Can happen!



Alternate paths:

$s \rightarrow u \rightarrow v \rightarrow t$

& $s \rightarrow v \rightarrow u \rightarrow t$

\Rightarrow Total runtime:
 $O(E \cdot V \cdot C)$

→ This running time makes F-F alg. be a pseudopolynomial algorithm = polynomial in sum of input numbers i.e., polynomial in unary encoding of input (NOT binary)

7

② If capacities rational
⇒ running time still finite (& pseudopoly.)

③ If capacities are real
⇒ potentially infinite runtime!

Important observations: (+)

→ If residual capacities are all integral
⇒ $C_f(p)$ is integral as well

→ If $C_f(p)$ is integral & f is integral (i.e., $f(u,v) \in \mathbb{Z} \forall (u,v)$)
⇒ residual capacities stay integral after augmentation
 & the flow

⇒ If network has all capacities integral then F-F alg. will find a max flow with all $f(u,v)$ being integral

⇒ Flow integrality thm: If $G=(V,E,s,t,c)$ has all capacities integral, then \exists a flow f s.t.
 $|f|=F^*$ & both F^* and f are integral

Note: There still can \exists max flow f' that is NOT integral (but $|f'|=F^*$ will still be integral)

Very useful when applying max flow

From now on: Assume capacities integral

(8)

\Rightarrow F-F runtime: $O(V \cdot EC)$

\rightarrow OK-ish when $C = \max_e c(e)$ small, e.g., $C=1$
but not so much otherwise
unit-capacity case

Can we get a better algorithm for large C ?

Idea: Ford-Fulkerson alg. just picks any augmenting path.
Is there a smart choice of augmenting path?

Maximum bottleneck path = augmenting path P that maximizes bottleneck capacity $c_P(P)$

Exercise: Max bottleneck path can be found in $O(E \log V)$ time
(Do binary search to find the max capacity C^*
Note: $C \leq C^*$ iff \exists s-t path in G_P after removing all edges (u,v) with $c_P(u,v) < C$)

Maximum bottleneck path (MBP) algorithm:

$\rightarrow f(u,v) = 0 \quad \forall_{u,v}$

\rightarrow While \exists augmenting path:

\rightarrow Find augmenting path P with maximum bottleneck capacity $c_P^* = c_P(P)$

\rightarrow Augment the flow with P
(so, $|f|$ increases by c_P^*)

\rightarrow Output f

What is the running time?

Claim: In any graph $G = (V, E, s, t, c)$, there exists an s - t path P in G with

$$c(P) = \min_{e \in P} c(e) \geq \frac{F^*}{m}$$

$$\begin{aligned} m &= |E| \\ m &= |V| \end{aligned}$$

9

Proof: Let f^* be the max flow in G

\Rightarrow By flow decomposition lemma, f^* decomposes into $\leq m$ flow paths & cycles

\Rightarrow The whole value $|f^*| = F^*$ of flow pushed by f^* is supported on $\leq m$ flow paths
(cycles do not push any flow from s to t)

\Rightarrow By averaging argument, at least one, P , of these $\leq m$ flow paths has to carry $\geq \frac{F^*}{m}$ flow

$$\Rightarrow c(P) \geq \frac{F^*}{m}$$

□

Corollary: MBP alg. runs in $O(m^2 \log n \log n C)$ time

Proof: \rightarrow Let F^* be the max flow value of the input network G

\rightarrow Let F_j be the amount of flow pushed by the MBP alg. up until iteration j

\rightarrow Let $F_j^* = F^* - F_j$ be the amount of flow "missing" at the beginning of iteration j
(= flow still to be pushed to get max flow)

\Rightarrow Max flow value of the residual graph G_f in j -th iteration is F_j^*

⇒ By applying the Claim to G_f we get that the (residual) bottleneck capacity of the max bottleneck capacity augmenting path P we found has to be

$$c_f(P) \geq \frac{F_j^*}{m}$$

$$\Rightarrow \forall_j F_{j+1}^* \leq \left(1 - \frac{1}{m}\right) F_j^* \leq \left(1 - \frac{1}{m}\right)^j F_1^* = \left(1 - \frac{1}{m}\right)^j F^*$$

⇒ Set $j^* \leftarrow m \ln(F^*) = O(m \log(nc))$ (since $F^* \leq nc$)
After j^*+1 iterations, we have

$$F_{j^*+2}^* \leq \left(1 - \frac{1}{m}\right)^{j^*+1} F^* < \left(1 - \frac{1}{m}\right)^{j^*} F^* \leq 1$$

⇒ But, by Flow integrality then,

if $F_{j^*+2}^* < 1$ then $F_{j^*+2}^* = 0$

Recall:
 $\left(1 - \frac{1}{k}\right)^k \leq e^{-1}$

⇒ MBP alg. terminates after $j^*+1 = O(m \log nc)$ iterations

⇒ Total running time: $O(m^2 \log n \cdot \log nc)$ [3]

What did we gain?

→ When C small, e.g., $C = o(n)$, vanilla F-F alg. still better
 $O(mn)$ vs. $O(m^2 \log^2 n)$

Note: For connected graphs, $m \geq n-1$

→ But, the runtime dependence only logarithmic in C

⇒ MBP is a (weakly) polynomial time alg. !

Runtime polynomial in
binary represent. of numbers


→ Runtime that is polynomial in $\log C$ is usually sufficient... but can we have NO dependency on C ? (11)

Yes! Again: use a "smart" choice (= strongly polynomial runtime) of augmenting path in FF alg.

Edmonds-Karp alg.: Variant of F-F alg. in which we always choose the augmenting path with fewest number of edges (i.e., shortest path if each edge in G_f has length 1) [1972]

Note: Can find such path in $O(n)$ time via BFS exploration

This choice avoids the "bad" instance for F-F alg.



E-K alg. (& MBP alg.!) need only two iterations to find max flow here

Claim: Edmonds-Karp alg. runs in $O(m^2 n)$ time

Proof: See appendix of recitations notes

(You are NOT responsible for this proof)

↖ Indeed, no dependency on C !

Highlights of later work on max flow algorithms:

→ Strongly polynomial time:

- [King Rao Tarjan '94]: $O(mn \log \frac{m}{n \log n})$

(= $O(mn)$ if $m = n^{1+\epsilon}$ $\epsilon > 0$)

- [Orlin '13]: $O(mn)$

→ (Weakly) polynomial time:

- [Goldberg Rao '98]: $O(\min\{m^{\frac{2}{3}}, mn^{\frac{2}{3}}\} \log \frac{m^2}{n} \log C)$


- [Lee Sidford '14]: $O(m\sqrt{n} \log^{\alpha(n)} n)$

- unit-capacity: [Madry '13]: $O(m^{\frac{10}{7}} \log^{\alpha(n)} n)$

Applications of Max Flow

(12)

Bipartite graph = graph $G=(V,E)$ with two sides L and R s.t. all edges go between L and R



In other words: $G=(V,E)$ s.t. $V=L \cup R$ & $E \subseteq L \times R$

↑
disjoint
union

↑
Cartesian
product

Matching = subset of edges s.t. no two of them share a vertex

Maximum bipartite matching problem: Given bipartite $G=(L \cup R, E)$ find a matching of maximum cardinality

Think: L = clients / jobs / groups / ...

R = servers / machines / slots / ... exclusive server / machine / ...


(Here: each client/job/... needs exactly one

Matching = assignment of clients to servers, jobs to machines, groups to slots...

Maximum matching = assignment that satisfies the most clients / jobs / groups

clients

servers



maximum
matching

edge (u,v) iff client u can be served by server v

Matching ensures that:

- each client uses ≤ 1 server
- each server serves ≤ 1 clients

How to compute maximum bipartite matching?

Reduce it to max flow problem!


We already know how to do that

= transform (quickly) an instance of max matching problem to an instance of max flow problem so as solving max flow instance gives you a solution to the max matching instance

Reduction:

(Note: Can be done in $O(m+n)$ time)

13



- Direct all the edges from L to R
- Add source s & sink t
- Add edges from s to each $u \in L$ & from each $v \in R$ to t
- Set all capacities to 1

Note: All capacities in G' are = 1 (and thus integral) $\leftarrow C=1$
 \Rightarrow By flow integrality then, integral max flow f^* exists
 \Rightarrow Ford-Fulkerson alg. will find such f^* in $O(mn)$ time

Claim: $|f^*| = |M^*|$, where $M^* = \text{max matching in } G$

Proof: $|f^*| \geq |M^*|$: Let $M^* = \{(u_1, v_1), (u_2, v_2), \dots, (u_k, v_k)\}$

A flow f that pushes unit flow on each path

$\langle s, u_i, v_i, t \rangle$ has a value $|f| = k = |M^*|$

But, obviously, $|f| \leq |f^*| = F^*$

$|f^*| \leq |M^*|$: f^* is integral, so is supported on exactly $|f^*|$ edges in $L \times R$. By feasibility & flow conserv.
none of these edges share an endpoint

\Rightarrow they form a matching of size $|f^*|$

$\Rightarrow |f^*| \leq |M^*|$

13

$\{s\} \cup L$ is an s - t cut containing only edges in $L \times R$

$f^*(\{s\} \cup L) = |f^*|$

Max bipartite matching alg.

- Compute G' $O(m+n)$ time
- Run F-F alg. on G' to compute f^* $O(mn)$ time
- Return a matching of edges in $L \times R$ that support (integral) f^* $O(m+n)$ time

Runtime: $O(mn)$ Can we do better?

[Hopcroft Karp '73]

$O(m\sqrt{n})$

[Mucha Samkowski '04]

$O(n^{\omega})$

[Madry '13]

$O(m^{10/9} \log^{\alpha(n)} n)$

exponent of fast matrix mult. ≤ 2.376 goal in dense graphs
($m = \Omega(n^2)$)goal in sparse graphs
($m = O(n)$)Another application of max flow:

Baseball elimination → see next page